



---

# CS5011 A4

---

200010781



MAY 10, 2022

## Table of Contents

|  |           |
|--|-----------|
| <b>1 Introduction .....</b>              | <b>2</b>  |
| <b>2 Design and Implementation .....</b> | <b>2</b>  |
| 2.1 Problem Overview .....               | 2         |
| 2.2 System Architecture .....            | 3         |
| 2.3 Implementation .....                 | 3         |
| 2.3.1 Part 1 .....                       | 3         |
| 2.3.2 Part 2 .....                       | 4         |
| 2.3.3 Part 3 .....                       | 4         |
| <b>3 Test Summary .....</b>              | <b>5</b>  |
| <b>4 Evaluation .....</b>                | <b>8</b>  |
| <b>5 Conclusion .....</b>                | <b>12</b> |

# 1 Introduction

This report describes the fourth assignment for CS5011 module. General information can be seen in the Table 1 below.

Table 1 General information of the assignment

|                           |              |
|---------------------------|--------------|
| <b>Assignment Number</b>  | Assignment 4 |
| <b>Date of Submission</b> | 10 May 2022  |
| <b>Student ID</b>         | 200010781    |
| <b>Word Count</b>         | 1998         |

There are three parts in this assignment, including building an ANN question classifier for the TREC dataset in part 1, implementing an Embedding layer to deal with the issue of slow matrix computation in part 2, using pre-trained word embeddings for setting ANN weights in part 3. The completion of each agent is described as Table 2 below.

Table 2 Completion of each part

|        |                 |                          |
|--------|-----------------|--------------------------|
| Part 1 | Build ANN       | attempted, fully working |
| Part 2 | Embedding layer | attempted, fully working |
| Part 3 | Initial weights | attempted, fully working |

About running the source code, several steps are needed.

1. Enter the terminal.
2. Change to the CS5011\_A4 directory.
3. Compile all source code
4. Run the source code

```
java A4Main <part1/part2/part3> <seed> <trainFile> <devFile> <testFile>  
<vocabFile> <classesFile>
```

## 2 Design and Implementation

### 2.1 Problem Overview

This assignment focuses on constructing and use Artificial Neural Networks (ANNs) to build a question classifier. In this assignment, the PEAS model is described as Table 3 below.

Table 3 PEAS model for assignment 2

| Agent                     | Performance Measure              | Environment                | Actuator           | Sensor    |
|---------------------------|----------------------------------|----------------------------|--------------------|-----------|
| Question classifier agent | Training time, Classify accuracy | Artificial Neural Networks | Classify questions | Using ANN |

## 2.2 System Architecture

All source code can be found in the src directory. The main function is addressed in A4Main.java file which locates in src. In order to keep the structure of this project, there are no packages in the src directory. There are 7 classes related to the functionalities in this assignment.

- A4Main: This class includes the main function of the assignment.
- A4Dataset: This class inherit Dataset class. It mainly aims at reading different datasets. `fromFile()` method is used to load training, validation and testing datasets and transfer them into one-hot-encoding format. `readVocab()` method is used to read vocabulary file in part 3 and store the provided words and pre-trained vector values for each word.
- EmbeddingBag: A class for Embedding bag layers. There are `forward()` and `backward()` methods for forward and backward propagations. `getX()` method is used to get the word indices from input object, which will be used in the backward process. `initW()` method aims at initializing the weight as required in part 3.
- Agent: This class includes the `train()`, `eval()` and `fromBatch()` methods to support training and evaluation processes in agents.
- P1: This class inherit Dataset class. It mainly aims at part 1 of the assignment.
- P2: This class inherit Dataset class. It mainly aims at part 2 of the assignment.
- P3: This class inherit Dataset class. It mainly aims at part 3 of the assignment.

## 2.3 Implementation

### 2.3.1 Part 1

This part focuses on building an ANN question classifier for the TREC dataset.

At first, I create three objects of A4Dataset named trainset, devset and testset to represent the training, validation, and testing datasets with batch size set. Then, it calls `fromFile()` method in A4Dataset to load datasets and change data into one-hot encoding format.

A network is constructed as required in the constructor of P1 class. There

are three hidden layers. The size of the first hidden layer is 100, and the sizes for other two are both 200. The outdims equals to 50 since there are 50 classes. In this assignment, I use cross entropy to calculate loss and stochastic gradient descent for optimization.

After building the network successfully, I call `train()` method for training. Lastly, I call `eval()` to perform the ANN on test set.

### 2.3.2 Part 2

This part focuses on using an Embedding layer to deal with the issue of slow matrix computation. P2 class focuses on this part.

General process is similar as part 1. The main difference is that part 2 uses `EmbeddingBag` to represent the embedding layer in your ANN model.

In the `forward()` method in `EmbeddingBag` class, I get the words indices of each sample from input object through `getX()` method at first. For each sample, I get the indices of each word and use it to find the corresponding weight. Then, sum up all weights and put the sum into Y matrix.

In the `backward()` method in `EmbeddingBag` class, I get the words indices of each sample and store it in variable `num`. Then, for each output node, I sum up `gy` and `gw`, where `gy` refers to the gradient of output Y and `gw` refers to the gradient of weight W, and this sum result will be put into `gW` matrix.

Similar as part 1, after building the network successfully, I call `train()` method for training. Lastly, I call `eval()` to perform the ANN on test set.

### 2.3.3 Part 3

This part focuses on using pre-trained word vector from GloVe to initialize weights in embedding layer. P3 class focuses on this part.

General process is similar as part 2. The main difference is that P2 uses random initialization but P3 uses pre-trained word vector to set weights.

At the beginning, I create a new `A4Dataset` named `vocabset`, and load it from input path through `readVocab()` method in `A4Dataset`. This method put all words into an `ArrayList<String>` named `word` and corresponding vector into an `ArrayList<double[]>` named `vectorValues`.

A new constructor is created in `EmbeddingBag` class to pass the `vocabset`. `initW()` method is used to initialize the weight using pre-trained vectors. In this method, I use `put()` method of `DoubleMatrix` to set the value of each weight.

Similar as part 2, after building the network successfully, I call `train()` method for training. Lastly, I call `eval()` to perform the ANN on test set.

### 3 Test Summary

This section mainly discusses the correctness of the system.

- Part 1

I show screen shots of the starting and ending epochs of part 1 using starting hyper-parameters with 123 seed. At the beginning, it prints the number of samples in datasets and the network we used in ANN. The training process will stop when there is no improvement for 10 epochs. At the end, I print the accuracies of three datasets and total running time.

```
Loading data...
train: 5000 instances
dev: 452 instances
test: 500 instances

Creating network...
(
  Linear: 3250 in, 100 out
  ReLU
  Linear: 100 in, 200 out
  ReLU
  Linear: 200 in, 200 out
  ReLU
  Linear: 200 in, 50 out
  Softmax
)

Training...
epoch: 0    loss: 333.0530  train-accuracy: 0.1996  dev-accuracy: 0.1770
epoch: 1    loss: 271.8985  train-accuracy: 0.3576  dev-accuracy: 0.3473
epoch: 2    loss: 238.2388  train-accuracy: 0.4658  dev-accuracy: 0.4425
epoch: 3    loss: 213.7707  train-accuracy: 0.5260  dev-accuracy: 0.4889
epoch: 4    loss: 191.0124  train-accuracy: 0.5654  dev-accuracy: 0.5066
```

Figure 1 The start of part 1 using starting hyper-parameters with 123 seed

```
not at peak 6 times consecutively
epoch: 31    loss: 6.8429  train-accuracy: 0.9932  dev-accuracy: 0.7434
not at peak 7 times consecutively
epoch: 32    loss: 5.6953  train-accuracy: 0.9928  dev-accuracy: 0.7500
not at peak 8 times consecutively
epoch: 33    loss: 5.0977  train-accuracy: 0.9930  dev-accuracy: 0.7633
epoch: 34    loss: 4.6926  train-accuracy: 0.9750  dev-accuracy: 0.7389
not at peak 1 times consecutively
epoch: 35    loss: 4.6368  train-accuracy: 0.9942  dev-accuracy: 0.7456
not at peak 2 times consecutively
epoch: 36    loss: 4.0800  train-accuracy: 0.9932  dev-accuracy: 0.7456
not at peak 3 times consecutively
epoch: 37    loss: 3.8159  train-accuracy: 0.9950  dev-accuracy: 0.7500
not at peak 4 times consecutively
epoch: 38    loss: 3.4911  train-accuracy: 0.9952  dev-accuracy: 0.7456
not at peak 5 times consecutively
epoch: 39    loss: 3.7010  train-accuracy: 0.9950  dev-accuracy: 0.7611
not at peak 6 times consecutively
epoch: 40    loss: 3.7738  train-accuracy: 0.9946  dev-accuracy: 0.7611
not at peak 7 times consecutively
epoch: 41    loss: 3.4299  train-accuracy: 0.9950  dev-accuracy: 0.7522
not at peak 8 times consecutively
epoch: 42    loss: 3.3176  train-accuracy: 0.9952  dev-accuracy: 0.7566
not at peak 9 times consecutively
epoch: 43    loss: 3.5637  train-accuracy: 0.9956  dev-accuracy: 0.7500
not at peak 10 times consecutively

training is finished

Test accuracy: 0.7620
Total running time: 337518ms
```

截图(Alt + /)

Figure 2 The end of part 1 using starting hyper-parameters with 123 seed

- Part 2

In order to make sure my implementation is correct, I use gradient checker and compute manually to make sure the correctness of part 2.

```

--- Test Classification ---
(
  Embedding: 5 rows, 10 dims
  Sigmoid
  Linear: 10 in, 20 out
  ReLU
  Linear: 20 in, 6 out
  Softmax
)
CrossEntropyLoss
correct backward for weights

```

Figure 3 Use GradientChecker to ensure the correctness

In figure 4 and figure 5, I show screen shots of the starting and ending epochs of part 2 using starting hyper-parameters with 123 seed. Similar as part 1, it prints the number of samples in datasets and the network we used in ANN at the beginning. The training process will stop when there is no improvement for 10 epochs. At the end, I print the accuracies of three datasets and total running time as well. In part 2, I fix the issue of slow matrix computation. Therefore, the total training time is far less than part 1.

```

Loading data...
train: 5000 instances
dev: 452 instances
test: 500 instances

Creating network...
(
  Embedding: 3250 rows, 100 dims
  ReLU
  Linear: 100 in, 200 out
  ReLU
  Linear: 200 in, 200 out
  ReLU
  Linear: 200 in, 50 out
  Softmax
)

Training...
epoch:  0      loss: 335.5717  train-accuracy: 0.1826  dev-accuracy: 0.1615
epoch:  1      loss: 275.0595  train-accuracy: 0.3388  dev-accuracy: 0.3252
epoch:  2      loss: 240.1332  train-accuracy: 0.4722  dev-accuracy: 0.4535

```

Figure 4 The start of part 2 using starting hyper-parameters with 123 seed

```

epoch: 27    loss: 24.2017    train-accuracy: 0.9864    dev-accuracy: 0.7412
epoch: 28    loss: 8.1618    train-accuracy: 0.9922    dev-accuracy: 0.7434
epoch: 29    loss: 6.5803    train-accuracy: 0.9930    dev-accuracy: 0.7124
not at peak 1 times consecutively
epoch: 30    loss: 6.9788    train-accuracy: 0.9932    dev-accuracy: 0.7323
not at peak 2 times consecutively
epoch: 31    loss: 6.6636    train-accuracy: 0.9930    dev-accuracy: 0.7345
not at peak 3 times consecutively
epoch: 32    loss: 5.3505    train-accuracy: 0.9940    dev-accuracy: 0.7389
not at peak 4 times consecutively
epoch: 33    loss: 4.7187    train-accuracy: 0.9930    dev-accuracy: 0.7323
not at peak 5 times consecutively
epoch: 34    loss: 4.4634    train-accuracy: 0.9838    dev-accuracy: 0.7212
not at peak 6 times consecutively
epoch: 35    loss: 4.4228    train-accuracy: 0.9950    dev-accuracy: 0.7345
not at peak 7 times consecutively
epoch: 36    loss: 3.9739    train-accuracy: 0.9936    dev-accuracy: 0.7412
not at peak 8 times consecutively
epoch: 37    loss: 3.5922    train-accuracy: 0.9942    dev-accuracy: 0.7367
not at peak 9 times consecutively
epoch: 38    loss: 3.4172    train-accuracy: 0.9950    dev-accuracy: 0.7257
not at peak 10 times consecutively

training is finished

Test accuracy: 0.7560
Total running time: 63842ms

```

Figure 5 The end of part 2 using starting hyper-parameters with 123 seed

### ● Part 3

In figure 6 and figure 7, I show screen shots of the starting and ending epochs of part 3 using starting hyper-parameters with 123 seed. Similar as previous parts, it prints the number of samples in datasets and the network we used in ANN at the beginning. Here, we use data from part3 directory, so the number of nodes in input layer is 8595. At the end, I print the accuracies of three datasets and total running time as well.

```

Loading data...
train: 5000 instances
dev: 452 instances
test: 500 instances

Creating network...
(
  Embedding: 8595 rows, 100 dims
  ReLU
  Linear: 100 in, 200 out
  ReLU
  Linear: 200 in, 200 out
  ReLU
  Linear: 200 in, 50 out
  Softmax
)

Training...
epoch: 0    loss: 318.5388    train-accuracy: 0.2790    dev-accuracy: 0.2699
epoch: 1    loss: 266.7840    train-accuracy: 0.4110    dev-accuracy: 0.3894
epoch: 2    loss: 236.8816    train-accuracy: 0.4308    dev-accuracy: 0.4137

```

Figure 6 The start of part 3 using starting hyper-parameters with 123 seed



```

not at peak 5 times consecutively
epoch: 38    loss: 3.8142    train-accuracy: 0.9982    dev-accuracy: 0.6858
epoch: 39    loss: 3.2459    train-accuracy: 0.9982    dev-accuracy: 0.6814
not at peak 1 times consecutively
epoch: 40    loss: 3.1095    train-accuracy: 0.9984    dev-accuracy: 0.6770
not at peak 2 times consecutively
epoch: 41    loss: 3.9121    train-accuracy: 0.9968    dev-accuracy: 0.6659
not at peak 3 times consecutively
epoch: 42    loss: 3.0574    train-accuracy: 0.9988    dev-accuracy: 0.6748
not at peak 4 times consecutively
epoch: 43    loss: 2.6489    train-accuracy: 0.9992    dev-accuracy: 0.6792
not at peak 5 times consecutively
epoch: 44    loss: 2.2214    train-accuracy: 0.9968    dev-accuracy: 0.6792
not at peak 6 times consecutively
epoch: 45    loss: 1.7818    train-accuracy: 0.9976    dev-accuracy: 0.6858
not at peak 7 times consecutively
epoch: 46    loss: 2.1613    train-accuracy: 0.9990    dev-accuracy: 0.6836
not at peak 8 times consecutively
epoch: 47    loss: 1.8528    train-accuracy: 0.9984    dev-accuracy: 0.6858
not at peak 9 times consecutively
epoch: 48    loss: 1.7184    train-accuracy: 0.9988    dev-accuracy: 0.6615
not at peak 10 times consecutively

training is finished

Test accuracy: 0.7300
Total running time: 107545ms

```

Figure 7 The end of part 2 using starting hyper-parameters with 123 seed

## 4 Evaluation

This section mainly discusses the performance of different parts. Firstly, I use different hyper-parameters to compare their effects to performance in part 1. Secondly, I use 6 different seeds to compare how seed effects the performance. Finally, I calculate the average epochs, loss, running time and accuracy rate of three parts to compare their performance.

### 1. Compare the effects of hyper-parameters to performance in part 1

Using the starting hyper-parameters: learning rate: 0.1; maximum number of epochs: 500; patience parameter: 10; batch size: 50, the performance of part 1 is shown in the table below.

Table 4 Performance of part 1 using starting hyper-parameters with 123 seed

| Seed (123)          | Num of epochs | Loss   | Running time(ms) | Training Accuracy(%) | Validation Accuracy(%) | Testing Accuracy(%) |
|---------------------|---------------|--------|------------------|----------------------|------------------------|---------------------|
| Starting Parameters | 43            | 3.5637 | 323257           | 99.56                | 75.00                  | 76.20               |

In order to find out the effects of each hyper-parameter to performance in part 1, I use Variable-controlling approach. I change each hyper-parameter one by one and record the performance data.

Table 5 Performance of part 1 using different parameters with 123 seed

| Seed (123)     | Num of epochs | Loss      | Running time(ms) | Training Accuracy(%) | Validation Accuracy(%) | Testing Accuracy(%) |
|----------------|---------------|-----------|------------------|----------------------|------------------------|---------------------|
| nEpoch=20      | 19            | 27.3746   | 145406           | 95.28                | 72.57                  | 72.6                |
| nEpoch=600     | 43            | 3.5637    | 379002           | 99.56                | 75.00                  | 76.2                |
| Patience=5     | 29            | 12.9134   | 224711           | 99.12                | 73.45                  | 75.8                |
| Patience=20    | 53            | 2.5983    | 420757           | 99.50                | 74.34                  | 75.4                |
| Batchsize=20   | 39            | 5.1523    | 308500           | 99.60                | 74.78                  | 79.8                |
| Batchsize=100  | 53            | 3.3467    | 426354           | 99.06                | 74.34                  | 74.6                |
| LearnRate=0.05 | 48            | 7.8978    | 362979           | 99.04                | 74.78                  | 74.2                |
| LearnRate=0.5  | 38            | 2.0997    | 276911           | 99.54                | 70.35                  | 77.2                |
| LearnRate=5    | 10            | 1485.4437 | 70589            | 7.84                 | 6.42                   | 24.6                |

### Analysis

Comparing Table 4 and Table 5, different parameters have different effects.

#### ■ nEpoch

Number of epochs controls the maximum iteration of the training process. As shown in Table 4, when setting nEpoch=500, we need 43 epochs to converge. When changing nEpoch to 600, there is no significant differences to performance. This is because the number of epochs we need is far more less than 500. However, if changing nEpoch to 20 as shown in the first line of Table 5, the number of epochs is exactly 20 (starting from 0). It is clear that the loss is higher than using large nEpoch and the accuracies of three datasets decrease. This is because we limit the maximum iteration to 20. The training process will stop when it reaches to 20 epochs where it has not reached the optimal solution yet.

#### ■ Patience

This parameter is used to control the number of epochs on which we would like to see no improvement. When setting patience=10, that means if we see no improvements for 10 epochs, the training process will stop. As shown in Table 5, if I set patience = 5, the training process will stop when no improvements for 5 epochs. In this situation, the loss

is higher than patience = 10 and the accuracies of three datasets are slightly lower. When setting patience = 20, the training process will stop when no improvements for 20 epochs. In this situation, the loss is slightly lower than patience = 10 and the accuracies of the three datasets are quite similar. According to these different results, smaller patience value may lead to a situation that model has not reached optimum and stops.

- **Batch size**

The batch size defines the number of samples that will be propagated through the network. When setting batchsize=20, networks train faster with less running time, but less accurate than using batchsize=50. When setting Batchsize = 100, time for training rises, but reaches higher accurate with smaller loss.

- **Learning rate**

The learning rate is a hyperparameter that controls how much to change the model in response to the estimated error each time the model weights are updated. It helps to minimize the network's loss function. When learning rate is small, such as learnRate=0.05, the networks spend longer time to train. This is because we need more steps to reach the optimum solution. However, if setting learnRate too large, such as learnRate=5, the network may not find the best solution, which leads to low accuracies for all datasets as shown in Table 5.

## **2. Compare the performance using different seeds**

In order to find out the effects for all three parts using different seeds, I use 6 different seed (123, 10, 50, 500, 5000, 10000) and record the number of epochs, loss, running time, and training accuracies for three datasets. All results are shown in the Table 6 below.

From Table 6, we can't see the obvious rules between seed and the performance of ANN. But different seeds do lead to different results.

Table 6 Results of three parts using different random seed

| Part   | Seed  | Num of epochs | Loss    | Running time(ms) | Training Accuracy(%) | Validation Accuracy(%) | Testing Accuracy(%) |
|--------|-------|---------------|---------|------------------|----------------------|------------------------|---------------------|
| Part 1 | 10    | 28            | 15.0994 | 297766           | 99.58                | 74.34                  | 79.6                |
| Part 1 | 50    | 33            | 5.1418  | 258132           | 99.32                | 73.23                  | 77.2                |
| Part 1 | 123   | 43            | 3.5637  | 323257           | 99.56                | 75.00                  | 76.2                |
| Part 1 | 500   | 54            | 2.6087  | 414731           | 99.56                | 73.67                  | 79.4                |
| Part 1 | 5000  | 48            | 3.8767  | 373886           | 99.52                | 73.67                  | 76.6                |
| Part 1 | 10000 | 49            | 6.4439  | 514828           | 99.60                | 74.34                  | 79.6                |
| Part 2 | 10    | 20            | 13.5113 | 46733            | 99.30                | 73.23                  | 77.2                |
| Part 2 | 50    | 28            | 18.1244 | 63197            | 99.60                | 73.01                  | 81.0                |
| Part 2 | 123   | 22            | 14.4975 | 51190            | 99.58                | 73.45                  | 77.8                |
| Part 2 | 500   | 25            | 11.1560 | 57140            | 99.56                | 72.35                  | 79.6                |
| Part 2 | 5000  | 37            | 6.9639  | 82312            | 99.62                | 75.22                  | 77.6                |
| Part 2 | 10000 | 33            | 7.3886  | 74215            | 99.62                | 73.01                  | 78.6                |
| Part 3 | 10    | 31            | 34.1079 | 119422           | 96.36                | 66.81                  | 71.8                |
| Part 3 | 50    | 32            | 18.7679 | 119031           | 99.02                | 69.03                  | 74.2                |
| Part 3 | 123   | 40            | 1.6524  | 146240           | 99.92                | 68.36                  | 76.0                |
| Part 3 | 500   | 30            | 4.1907  | 109854           | 99.84                | 67.26                  | 74.2                |
| Part 3 | 5000  | 42            | 1.5148  | 152376           | 99.92                | 67.26                  | 73.0                |
| Part 3 | 10000 | 53            | 1.2122  | 191365           | 99.92                | 66.15                  | 74.8                |

### 3. Compare the performance of three parts

In order to compare the performances of all three parts, I get the average results of using 6 different seeds. The results are shown in the Table 7 below.

Due to sparse input representation in part 1, it spends the longest time to train. In part 2, after fixing the issue of slow matrix computation, the running time decreases significantly. The testing accuracy of part 2 is slightly lower than part1, I assume it is because we did not use bias in embedding layer. In part 3, I use pre-trained embedding on your ANN models. However, there is no improvement in performance comparing to part 2. Besides, the time spent in part 3 is slightly longer than part 2. I assume this is because part 3 spend some time on assigning vector values to weights in embedding layer.

| Average (on 6 seeds) | Num of epochs | Loss    | Running time(ms) | Training Accuracy(%) | Validation Accuracy(%) | Testing Accuracy(%) |
|----------------------|---------------|---------|------------------|----------------------|------------------------|---------------------|
| Part 1               | 42.5          | 6.1224  | 363767           | 99.52                | 74.04                  | 78.1                |
| Part 2               | 27.5          | 11.9403 | 62464            | 99.54                | 73.38                  | 77.6                |
| Part 3               | 38            | 10.2409 | 139714           | 99.16                | 67.47                  | 74.0                |

## 5 Conclusion

There are three parts in this assignment, including building an ANN question classifier for the TREC dataset in part 1, implementing an Embedding layer to deal with the issue of slow matrix computation in part 2, using pre-trained word embeddings for setting ANN weights in part 3.

In part 1, I compare the effects of each parameter to the performance of ANN model. It can be seen that network trains faster with mini batches but lower accuracy. Too small nEpoch may cause a network stop early where it has not reached the optimum solution. Similar to parameter patience, smaller patience value may lead to a situation that model has not reached optimum and stops. It is tricky to choose the right learning rate in machine learning problem. It depends on different model with different loss function. Large learning rate may lead the network miss the minimum loss and small learning rate increases the training time.

There are no obvious rules between seed and the performance of ANN. After calculating the average results of three parts, part 2 saves more time than part 1 during training. Part 3 needs longer time than part 2 but less time than part 1. The reason might be assigning vector values to weights spends some time. In terms of accuracy, part 2 has slightly lower accuracy than part 1. I assume it is because we consider no bias in the embedding layer. There are no improvements in part 3 comparing to part 2.